

NBA 4920/6921 Lecture 13

Ridge Regression Application

Murat Unal

10/14/2021

```
rm(list=ls())
options(digits = 3, scipen = 999)
library(tidyverse)
library(ISLR)
library(cowplot)
library(ggcorrplot)
library(stargazer)
library(corr)
library(lmtest)
library(sandwich)
library(MASS)
library(car)
library(jtools)
library(caret)
library(leaps)
library(future.apply)
library(glmnet)
Hitters <- ISLR::Hitters
Hitters <- na.omit(Hitters)
set.seed(2)
```

Ridge Regression

```
x=model.matrix(Salary~.,Hitters)[-1]  
y=Hitters$Salary
```

- ▶ The **model.matrix()** function is particularly useful for creating x ; not only does it produce a matrix corresponding to the 19 predictors but it also automatically transforms any qualitative variables into dummy variables.
- ▶ The latter property is important because **glmnet()** can only take numerical, quantitative inputs.

- ▶ The **glmnet()** engine has an `alpha` argument that determines what type of model is fit.
- ▶ If `alpha=0` then a ridge regression model is fit, and if `alpha=1` then a lasso model is fit. We first fit a ridge regression model.
- ▶ By default, the **glmnet()** engine standardizes the variables so that they are on the same scale.

- ▶ **glmnet()** will fit ridge models across a wide range of λ values by default, 100 λ values that are data derived
- ▶ We could also run the engine for a grid of 100 values ranging from high to low λ

```
grid=10^seq(10,-2,length=100)  
grid[1]
```

```
[1] 100000000000
```

```
grid[100]
```

```
[1] 0.01
```

- For regression problems set `family="gaussian"`, and `family="binomial"` for classification

```
ridge.mod=glmnet(x,y,alpha=0,lambda=grid,family="gaussian")  
names(ridge.mod)
```

```
[1] "a0"      "beta"    "df"      "dim"     "lambda"  
[7] "nulldev" "npasses" "jerr"    "offset"  "call"
```


- ▶ Associated with each value of λ is a vector of ridge regression coefficients, stored in a matrix that can be accessed by `coef()`.
- ▶ In this case, it is a **20×100** matrix, with 20 rows (one for each predictor, plus an intercept) and 100 columns (one for each value of λ).

```
dim(coef(ridge.mod))
```

```
[1] 20 100
```

- ▶ As λ grows larger, our ridge regression coefficient magnitudes are more constrained.

```
plot(ridge.mod, xvar = "lambda")
```




```
#Display 1st lambda value
```

```
ridge.mod$lambda[1]
```

```
[1] 100000000000
```

```
# Display coefficients associated with
```

```
# 1st lambda value
```

```
coef(ridge.mod)[,1]
```

(Intercept)	AtBat	Hits	RBI
535.92569352090	0.00000005443	0.00000019746	0.00000007143
RBI	Walks	Years	CRBI
0.00000035272	0.00000041513	0.00000169771	0.00000000000
CHmRun	CRuns	CRBI	CRF
0.00000012972	0.00000003451	0.00000003561	0.00000000000
DivisionW	PutOuts	Assists	Errors
-0.00000780726	0.00000002180	0.00000000356	-0.00000000000

- ▶ We can use the `predict()` function for a number of purposes. For instance, we can obtain the ridge regression coefficients for a new value of λ , say 50:
- ▶ If the desired λ value is not included in the initial fit, then `glmnet()` performs linear interpolation to make predictions for the desired λ value.
- ▶ Linear interpolation usually works fine, but we could change this by calling `exact=TRUE`. This way predictions are to be made at values of λ not included in the original fit and the model is refit before predictions are made.

- ▶ Let's check if $\lambda = 50$ is used in the original fit

```
any(ridge.mod$lambda==50)
```

```
[1] FALSE
```

- ▶ Let's see if there's any difference

```
coef.approx <- predict(ridge.mod, s = 50, type = "coefficients",  
                        exact = FALSE)  
coef.exact <- predict(ridge.mod, s = 50, type = "coefficients",  
                      exact = TRUE, x=x,y=y)  
  
cbind(coef.approx, coef.exact)
```

```
20 x 2 sparse Matrix of class "dgCMatrix"
```

	1	1
(Intercept)	48.76610	48.272
AtBat	-0.35810	-0.353
Hits	1.96936	1.951
HmRun	-1.27825	-1.290
Runs	1.14589	1.156

- ▶ Let's now split the samples into a training set and a test set in order to estimate the test error of ridge regression

```
train=sample(1:nrow(x), 0.7*nrow(x))
```

- ▶ Next we fit a ridge regression model on the training set, and evaluate its RMSE on the test set, using $\lambda = 4$.
- ▶ Note the use of the `predict()` function again.
- ▶ This time we get predictions for a test set, by replacing `type="coefficients"` with the `newx` argument

```
ridge.mod=glmnet(x[train,],y[train],alpha=0,lambda=grid)
ridge.pred.lambda4=predict(ridge.mod,s=4,newx=x[-train,])
rmse.ridge.lambda4 = sqrt(mean((ridge.pred.lambda4-y[-train])^2))
rmse.ridge.lambda4
```

[1] 296

- ▶ Note that if we had instead simply fit a model with just an intercept, we would have predicted each test observation using the mean of the training observations.
- ▶ In that case, we could compute the test set MSE like this:

```
sqrt(mean((mean(y[train])-y[-train])^2))
```

```
[1] 405
```

- ▶ We could also get the same result by fitting a ridge regression model with a very large value of λ . Note that $1e10$ means 10^{10}

```
ridge.pred.lambdabig=predict(ridge.mod,s=1e10,newx=x[-train  
rmse.ridge.lambdabig <- sqrt(mean((ridge.pred.lambdabig-y[-  
rmse.ridge.lambdabig
```

```
[1] 405
```

- ▶ So fitting a ridge regression model with $\lambda = 4$ leads to a much lower -train RMSE than fitting a model with just an intercept.
- ▶ Let's now check whether there is any benefit to performing ridge regression with $\lambda = 4$ instead of just performing least squares regression.
- ▶ Recall that least squares is simply ridge regression with $\lambda = 0$

```
ridge.pred.lambda0=predict(ridge.mod,s=0,exact = TRUE,  
                           x=x[train,],y=y[train],newx=x[-train,])  
rmse.ridge.lambda0 <- sqrt(mean((ridge.pred.lambda0-y[-train])  
rmse.ridge.lambda0
```

```
[1] 300
```

- ▶ It looks like we are indeed improving over regular least-squares!

Cross-validation

- ▶ Instead of arbitrarily choosing $\lambda = 4$, it would be better to use cross-validation to choose the tuning parameter λ .
- ▶ We can do this by writing our own function as we did before or we could just use the built-in cross-validation function, `cv.glmnet()`.
- ▶ By default, the function performs 10-fold cross-validation, though this can be changed using the argument `folds`.
- ▶ We can define the loss used for cross-validation using `type.measure`

```
cv.out=cv.glmnet(x[train,],y[train],alpha=0,nfold=10,  
                 type.measure="mse")
```


```
bestlam=cv.out$lambda.min
```

```
bestlam
```

```
[1] 249
```

- ▶ The first vertical dashed line gives the value of λ that gives minimum mean cross-validated error. The second is the λ that gives an MSE within one standard error of the smallest.

```
plot(cv.out)
```



- ▶ Finally, we obtain the coefficients from our fitted model using the value of λ chosen by cross-validation.
- ▶ As expected, none of the coefficients are zero—ridge regression does not perform variable selection!

```
predict(ridge.mod,type ="coefficients",s=bestlam,  
        exact=TRUE, x=x[train,], y=y[train])
```

20 x 1 sparse Matrix of class "dgCMatrix"

	1
(Intercept)	6.57462
AtBat	-0.00753
Hits	0.94106
HmRun	1.42067
Runs	1.21523
RBI	1.18983
Walks	2.16999
Years	0.43523
CAtBat	0.00963
CHits	0.05862

```
final.pred = predict(ridge.mod, newx=x[-train,], s=bestlam,  
                      exact=TRUE, x=x[train,], y=y[train])  
rmse.ridge.lambdabest <- sqrt(mean((y[-train]-final.pred)^2))  
rmse.ridge.lambdabest
```

```
[1] 292
```

- ▶ Let's compare all test RMSEs:

```
rmse.ridge.lambdabig
```

```
[1] 405
```

```
rmse.ridge.lambda4
```

```
[1] 296
```

```
rmse.ridge.lambda0
```

```
[1] 300
```

```
rmse.ridge.lambdabest
```

```
[1] 292
```

- ▶ Ridge regression indeed obtains the minimum test RMSE!

```
ridge.pred.lambdamedian=predict(ridge.mod,s=10098,newx=x[-t  
rmse.ridge.lambdamedian <- sqrt(mean((ridge.pred.lambdamed  
rmse.ridge.lambdamedian
```

```
[1] 358
```