

NBA 4920/6921 Lecture 15

Elastic Net Application

Murat Unal

10/21/2021

```
rm(list=ls())
options(digits = 3, scipen = 999)
library(tidyverse)
library(ISLR)
library(ggplot2)
library(cowplot)
library(ggcorrplot)
library(stargazer)
library(corr)
library(lmtest)
library(sandwich)
library(MASS)
library(car)
library(jtools)
library(caret)
library(leaps)
library(future.apply)
library(glmnet)
Hitters <- ISLR::Hitters
Hitters <- na.omit(Hitters)
```

► Previous test RMSEs:

```
RMSE <- matrix(NA, ncol = 1, nrow = 8)
rownames(RMSE) <- c("rmse.ridge.lambdabig",
"rmse.ridge.lambda4", "rmse.ridge.lambda0",
"rmse.ridge.lambdabest", "rmse.lasso.lambda1se",
"rmse.lasso.lambdabest", "rmse.elnet.lambda1se",
"rmse.elnet.lambdabest")
RMSE[1:6,1] <- c(405, 296, 300, 292, 334, 297)
RMSE
```

	[,1]
rmse.ridge.lambdabig	405
rmse.ridge.lambda4	296
rmse.ridge.lambda0	300
rmse.ridge.lambdabest	292
rmse.lasso.lambda1se	334
rmse.lasso.lambdabest	297
rmse.elnet.lambda1se	NA
rmse.elnet.lambdabest	NA

Elastic net

- ▶ Now, there are two parameters to tune: λ and α .
- ▶ The **glmnet** package allows to tune λ via cross-validation for a fixed α , but it does not support α -tuning.

- ▶ Let's write our own loop that does the tuning
- ▶ First, we create a common `fold_id`, which just allows us to apply the same CV folds to each model.
- ▶ We then create a tuning grid that searches across a range of α s from 0-1, and empty columns where we'll dump our model results into.

```
# maintain the same folds across all models  
fold_id <- sample(1:10, size = length(y[train]),  
                  replace=TRUE)
```

```
# search across a range of alphas  
tuning_grid <- data.frame(  
  alpha      = seq(0, 1, by = .1),  
  mse_min    = NA, mse_1se    = NA,  
  lambda_min = NA, lambda_1se = NA)
```

- Now we can iterate over each α value, apply a CV elastic net, and extract the minimum and one standard error MSE values and their respective λ values.

```
for(i in seq_along(tuning_grid$alpha) ) {  
  # fit CV model for each alpha value  
  fit <- cv.glmnet(x[train,], y[train],  
                  alpha = tuning_grid$alpha[i],  
                  foldid = fold_id)  
  
  # extract MSE and lambda values  
  tuning_grid$mse_min[i]      <- fit$cvm[fit$lambda==  
                                         fit$lambda.min]  
  tuning_grid$mse_1se[i]     <- fit$cvm[fit$lambda==  
                                         fit$lambda.1se]  
  
  tuning_grid$lambda_min[i] <- fit$lambda.min  
  tuning_grid$lambda_1se[i] <- fit$lambda.1se  
}
```

```
tuning_grid %>% arrange(mse_min)
```

alpha	mse_min	mse_1se	lambda_min	lambda_1se
0.0	129229	162541	24.31	4451
1.0	129734	161372	2.79	115
0.9	129775	161084	3.11	128
0.8	129932	160957	3.49	144
0.7	130133	161130	3.64	165
0.1	130311	164170	135.91	959
0.6	130320	161324	3.87	192
0.2	130418	162119	89.83	526
0.5	130565	161697	4.23	231
0.3	130710	163761	65.73	385
0.4	130749	162338	5.29	289

- ▶ Extract the optimum *alpha* and λ values

```
best.index <- which.min(tuning_grid$mse_min)
best.alpha <- tuning_grid[best.index , "alpha"]
best.lambda <- tuning_grid[best.index , "lambda_min"]
best.lambda.1se <- tuning_grid[best.index , "lambda_1se"]
```

```
best.alpha
```

```
[1] 0
```

```
best.lambda
```

```
[1] 24.3
```

```
best.lambda.1se
```

```
[1] 4451
```



```
RMSE[7:8,1]<-c(rmse.elnet.lambda1se,  
               rmse.elnet.lambdabest)
```

```
RMSE
```

```
               [,1]  
rmse.ridge.lambdabig    405  
rmse.ridge.lambda4      296  
rmse.ridge.lambda0      300  
rmse.ridge.lambdabest   292  
rmse.lasso.lambda1se    334  
rmse.lasso.lambdabest   297  
rmse.elnet.lambda1se    329  
rmse.elnet.lambdabest   295
```

- ▶ We could also use the **caret** package to do cross-validation for both α and λ
- ▶ The package has the `train()` meta engine (aggregator) that allows us to apply almost any direct engine with `method()`

```
cv_10 = trainControl(method = "cv", number = 10)

grid = expand.grid(alpha = seq(0,1,by=0.1),
                  lambda = 10^seq(3,-2,length=100))

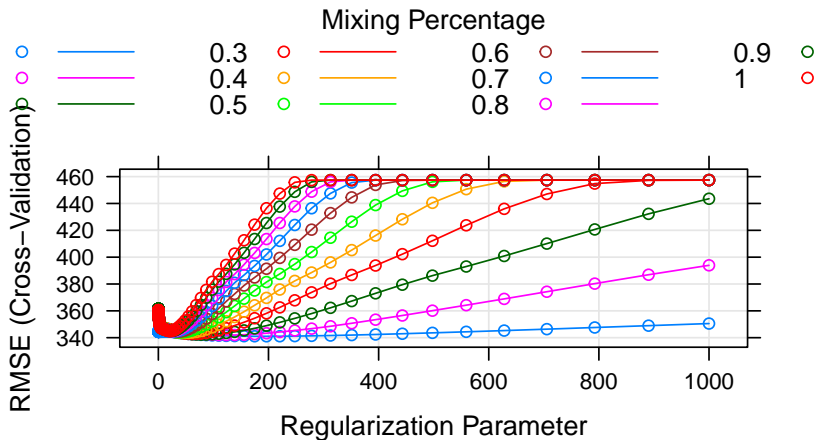
elnet = train(
  Salary ~ .,
  data = Hitters[train,],
  method = "glmnet",
  trControl = cv_10,
  preProcess = c("center", "scale"),
  tuneGrid = grid)
```

```
elnet$bestTune
```

	alpha	lambda
	86	0 196

```
alpha.best <- unlist(unname(elnet$bestTune[1]))  
lambda.best <- unlist(unname(elnet$bestTune[2]))
```

```
plot(elnet, xvar = "lambda")
```



- Final model with cross-validated parameters

```
elnet.final <- glmnet(x[train,],y[train],alpha=alpha.best)

elnet.final.lambdabest <- predict(elnet.final,
                                s=lambda.best,newx=x[-train,])

elnet.lambdabest.caret <- sqrt(mean((y[-train]-
                                elnet.final.lambdabest)^2 ))
elnet.lambdabest.caret
```

```
[1] 293
```

```
predict(elnet.final,s=lambda.best,type = "coefficients")[1]
```

(Intercept)	AtBat	Hits	HmRun	Runs
7.62142	-0.05566	1.02454	1.29116	1.24767
Walks	Years	CAtBat	CHits	CHmRun
2.31336	-0.31149	0.00929	0.06329	0.43651
CRBI	CWalks	LeagueN	DivisionW	PutOuts
0.11970	0.04249	17.35836	-135.45306	0.21112
Errors	NewLeagueN			
-1.94017	5.37944			