

NBA 4920/6921 Lecture 15

Elastic Net Application

Murat Unal

10/21/2021

```
rm(list=ls())
options(digits = 3, scipen = 999)
library(tidyverse)
library(ISLR)
library(ggplot2)
library(jtools)
library(caret)
library(leaps)
library(glmnet)
Hitters <- ISLR::Hitters
Hitters <- na.omit(Hitters)
set.seed(2)
```

- ▶ Let's create 2.order variables and interaction terms

```
x=model.matrix(Salary~.,Hitters)[-1]
y=Hitters$Salary

x.squared <- sapply(as.data.frame(x), function(i) i^2)

colnames(x.squared) <- paste0(colnames(x), "_sq", sep = "")

x.interact = model.matrix(~.^2,as.data.frame(x) )[,21:191]

x <- cbind(x, x.squared, x.interact)
rm(x.interact,x.squared)
dim(x)
```

```
[1] 263 209
```

```
train=sample(1:nrow(x), 0.7*nrow(x))  
ols.data <- as.data.frame(cbind(Salary=y[train],x[train,]))
```

► Matrix to store RMSEs:

```
RMSE <- matrix(NA, ncol = 1, nrow = 6)
rownames(RMSE) <- c("rmse.ridge.lambda0",
"rmse.ridge.lambdabest", "rmse.lasso.lambda1se",
"rmse.lasso.lambdabest", "rmse.elnet.lambda1se",
"rmse.elnet.lambdabest")
```

OLS

```
ols.mod <- lm(Salary~.,ols.data)
summ(ols.mod)
```

MODEL INFO:

Observations: 184

Dependent Variable: Salary

Type: OLS linear regression

MODEL FIT:

$F(183,0) = \text{NaN}$, $p = \text{NaN}$

$R^2 = 1.00$

Adj. $R^2 = \text{NaN}$

Standard errors: OLS

| | Est. | S.E. | t val. |
|-------------|---------|------|--------|
| (Intercept) | 1323.31 | | |

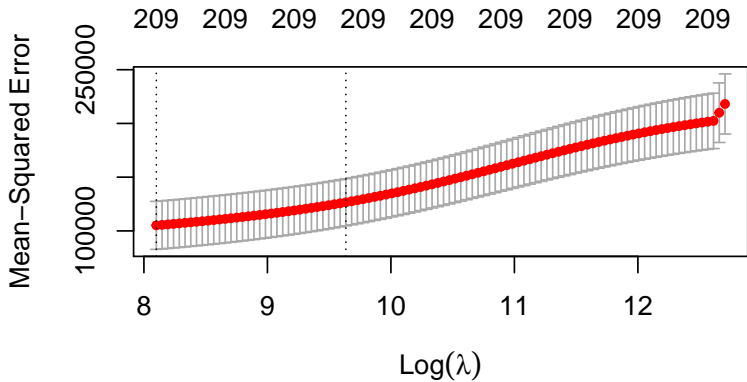
Ridge Regression

```
ridge.cv=cv.glmnet(x[train,],y[train]  
                  ,alpha=0,nfold=10,  
                  type.measure="mse")  
bestlam=ridge.cv$lambda.min  
bestlam
```

```
[1] 3296
```

► Plot

```
plot(ridge.cv)
```



► Predictions for $\lambda = 0$

```
ridge.pred.lambda0 = predict(ridge.cv,newx=x[-train,],  
                             s=0,exact=TRUE,  
                             x=x[train,],y=y[train])
```

```
rmse.ridge.lambda0 <- sqrt(mean((y[-train]-  
                                ridge.pred.lambda0)^2))
```

```
rmse.ridge.lambda0
```

```
[1] 556
```

► Predictions for best λ

```
ridge.pred.lambdabest = predict(ridge.cv,newx=x[-train,],  
                                s=bestlam,exact=TRUE,  
                                x=x[train,],y=y[train])
```

```
rmse.ridge.lambdabest <- sqrt(mean((y[-train]-  
                                ridge.pred.lambdabest)^2))
```

```
rmse.ridge.lambdabest
```

```
[1] 312
```

- ▶ Store the values in the RMSE matrix


```
RMSE[1:2,] <- c(rmse.ridge.lambda0,rmse.ridge.lambdabest)
```

Lasso Regression

```
lasso.cv=cv.glmnet(x[train,],y[train],alpha=1,  
                   nfold=10,type.measure="mse")
```

► Plot

```
plot(lasso.cv)
```



- ▶ Prediction using both values of λ , **lambda.min** and **lambda.1se**, the value of λ that gives the most regularized model such that the cross-validated error is within one standard error of the minimum.
- ▶ Save the outputs as **lasso.lambdabest** and **lasso.lambda1se**

```
lasso.pred.lambdabest=predict(lasso.cv,newx=x[-train,],
                             s=lasso.cv$lambda.min,exact=TRUE,
                             x=x[train,],y=y[train])
```

```
lasso.pred.lambda1se=predict(lasso.cv,newx=x[-train,],
                             s=lasso.cv$lambda.1se,exact=TRUE,
                             x=x[train,],y=y[train])
```

- Compute RMSE of the predictions

```
rmse.lasso.lambda.best <- sqrt(mean((y[-train]-  
                                     lasso.pred.lambda.best)^2))
```

```
rmse.lasso.lambda.1se <- sqrt(mean((y[-train]-  
                                     lasso.pred.lambda.1se)^2))
```

```
rmse.lasso.lambda.best
```

```
[1] 361
```

```
rmse.lasso.lambda.1se
```

```
[1] 287
```


► Store in RMSE

```
RMSE[3:4,] <- c(rmse.lasso.lambda1se,rmse.lasso.lambdabest)  
RMSE
```

| | [,1] |
|-----------------------|------|
| rmse.ridge.lambda0 | 556 |
| rmse.ridge.lambdabest | 312 |
| rmse.lasso.lambda1se | 287 |
| rmse.lasso.lambdabest | 361 |
| rmse.elnet.lambda1se | NA |
| rmse.elnet.lambdabest | NA |

Elastic net

- ▶ Now, there are two parameters to tune: λ and α .
- ▶ The **glmnet** package allows to tune λ via cross-validation for a fixed α , but it does not support α -tuning.

- ▶ Let's write our own loop that does the tuning
- ▶ First, we create a common `fold_id`, which just allows us to apply the same CV folds to each model.
- ▶ We then create a tuning grid that searches across a range of α s from 0-1, and empty columns where we'll dump our model results into.

```
# maintain the same folds across all models  
fold_id <- sample(1:10, size = length(y[train]),  
                  replace=TRUE)
```

```
# search across a range of alphas  
tuning_grid <- data.frame(  
  alpha      = seq(0, 1, by = .1),  
  mse_min    = NA, mse_1se    = NA,  
  lambda_min = NA, lambda_1se = NA)
```

- Now we can iterate over each α value, apply a CV elastic net, and extract the minimum and one standard error MSE values and their respective λ values.

```
for(i in seq_along(tuning_grid$alpha) ) {  
  # fit CV model for each alpha value  
  fit <- cv.glmnet(x[train,], y[train],  
                  alpha = tuning_grid$alpha[i],  
                  foldid = fold_id)  
  
  # extract MSE and lambda values  
  tuning_grid$mse_min[i]      <- fit$cvm[fit$lambda==  
                                         fit$lambda.min]  
  tuning_grid$mse_1se[i]      <- fit$cvm[fit$lambda==  
                                         fit$lambda.1se]  
  
  tuning_grid$lambda_min[i] <- fit$lambda.min  
  tuning_grid$lambda_1se[i] <- fit$lambda.1se  
}
```

```
tuning_grid %>% arrange(mse_min)
```

| alpha | mse_min | mse_1se | lambda_min | lambda_1se |
|-------|---------|---------|------------|------------|
| 0.7 | 90193 | 108536 | 9.03 | 101.4 |
| 0.8 | 90290 | 108465 | 6.87 | 88.8 |
| 1.0 | 90428 | 109678 | 11.57 | 77.9 |
| 0.6 | 90454 | 108687 | 10.54 | 118.3 |
| 0.5 | 90880 | 110071 | 13.87 | 148.8 |
| 0.9 | 90992 | 109461 | 6.11 | 82.6 |
| 0.4 | 91105 | 109946 | 18.17 | 177.5 |
| 0.3 | 91258 | 109891 | 23.12 | 225.9 |
| 0.2 | 91314 | 110403 | 34.69 | 323.5 |
| 0.1 | 91535 | 110824 | 66.22 | 562.7 |
| 0.0 | 105378 | 122038 | 3295.63 | 11571.6 |

- ▶ Extract the optimum *alpha* and λ values

```
best.index <- which.min(tuning_grid$mse_min)
best.alpha <- tuning_grid[best.index , "alpha"]
best.lambda <- tuning_grid[best.index , "lambda_min"]
best.lambda.1se <- tuning_grid[best.index , "lambda_1se"]
```

```
best.alpha
```

```
[1] 0.7
```

```
best.lambda
```

```
[1] 9.03
```

```
best.lambda.1se
```

```
[1] 101
```



```
RMSE[5:6,1]<-c(rmse.elnet.lambda1se,  
               rmse.elnet.lambdabest)
```

```
RMSE
```

```
               [,1]  
rmse.ridge.lambda0      556  
rmse.ridge.lambdabest   312  
rmse.lasso.lambda1se    287  
rmse.lasso.lambdabest   361  
rmse.elnet.lambda1se    295  
rmse.elnet.lambdabest   357
```