

NBA 4920/6921 Lecture 15

Tree Methods Application

Murat Unal

10/21/2021

```
rm(list=ls())
options(digits = 3, scipen = 999)
library(tidyverse)
library(ISLR)
library(cowplot)
library(ggcorrplot)
library(stargazer)
library(corr)
library(lmtest)
library(sandwich)
library(MASS)
library(car)
library(jtools)
library(caret)
library(leaps)
library(future.apply)
library(glmnet)
library(rpart)
library(rpart.plot)
library(ROCR)
```

```
Hitters <- ISLR::Hitters  
Hitters <- na.omit(Hitters)  
Carseats <- ISLR::Carseats  
Carseats = na.omit(Carseats)
```

```
train = sample(1:nrow(Hitters), 0.7*nrow(Hitters))
```

Regression Trees

- ▶ To train decision trees in R, we can use **caret** , which draws upon **rpart**
- ▶ You can get a list of models supported by **caret**,
`names(getModelInfo())`
- ▶ The tunable parameters for a given model,
`modelLookup("rpart")`
- ▶ To `train()` our model in **caret**
*Define the method as **rpart***
*The main tuning parameter is **cp**, the complexity parameter*

```

hit_tree = train(
  Salary ~ .,
  data = Hitters[train,],
  method = "rpart",
  trControl = trainControl("cv", number = 5),
  # tuneLength = 20
  tuneGrid = data.frame(cp = seq(0, 0.1, by = 0.001))
)
names(hit_tree)

```

[1]	"method"	"modelInfo"	"modelType"	"results"
[6]	"bestTune"	"call"	"dots"	"metric"
[11]	"finalModel"	"preProcess"	"trainingData"	"resample"
[16]	"perfNames"	"maximize"	"yLimits"	"times"
[21]	"terms"	"coefnames"	"contrasts"	"xlevels"

► To get the CV-chosen final tree

```
hit_tree$finalModel
```

```
n= 184
```

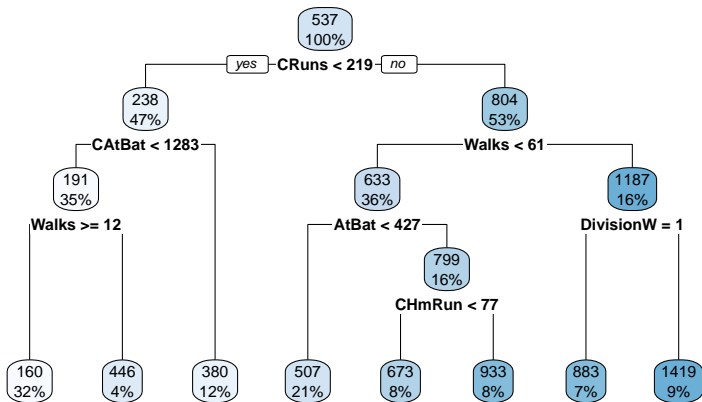
```
node), split, n, deviance, yval
```

```
* denotes terminal node
```

```
1) root 184 40400000 537
  2) CRuns< 218 87 5140000 238
    4) CAtBat< 1.28e+03 65 4150000 191
      8) Walks>=11.5 58 311000 160 *
      9) Walks< 11.5 7 3330000 446 *
    5) CAtBat>=1.28e+03 22 401000 380 *
  3) CRuns>=218 97 20500000 804
    6) Walks< 61 67 5700000 633
      12) AtBat< 426 38 1720000 507 *
      13) AtBat>=426 29 2570000 799
        26) CHmRun< 76.5 15 335000 673 *
        27) CHmRun>=76.5 14 1750000 933 *
```


- ▶ To plot the CV-chosen tree, we need to
 1. extract the fitted model, e.g., `hit_tree$finalModel`
 2. apply a plotting function e.g., `rpart.plot()` from **rpart.plot**


```
rpart.plot(hit_tree$finalModel)
```



- Plot the performance metric against the complexity parameter

```
cp.data = data.frame("cp" = hit_tree$results[[1]],  
                     "RMSE" = hit_tree$results[[2]])  
ggplot(cp.data, aes(x=cp,y=RMSE))+  
  geom_point()+  
  geom_line()
```



- ▶ Make predictions on the test data

```
pred.tree <- predict(hit_tree, Hitters[-train,])  
mean((Hitters[-train, "Salary"] - pred.tree)^2)
```

```
[1] 95393
```

Classification Trees

- ▶ We'll use the Carseats data set from ISLR

```
dim(Carseats)
```

```
[1] 400  11
```

```
names(Carseats)
```

```
[1] "Sales"           "CompPrice"       "Income"          "Advertising"
[6] "Price"           "ShelveLoc"      "Age"             "Education"
[11] "US"
```

- ▶ Let's modify the response from its original numeric variable to a categorical variable with two levels: high and low

```
Carseats$Sales = as.factor(ifelse(Carseats$Sales <= 8, "Low", "High"))
```

```
train = sample(1:nrow(Carseats), 0.7*nrow(Carseats))
```

- ▶ We change the performance metric to ROC, which is simply the AUC.

```
sales_tree = train(
  Sales ~ .,
  data = Carseats[train,],
  method = "rpart",
  trControl = trainControl("repeatedcv", number = 10,
                           repeats = 3,
                           summaryFunction = twoClassSummary,
                           classProbs = TRUE),
  metric = "ROC",
  # tuneLength=20
  tuneGrid = data.frame(cp = seq(0, 0.2, by = 0.005))
)
```

sales_tree

CART

280 samples

10 predictor

2 classes: 'High', 'Low'

No pre-processing

Resampling: Cross-Validated (10 fold, repeated 3 times)

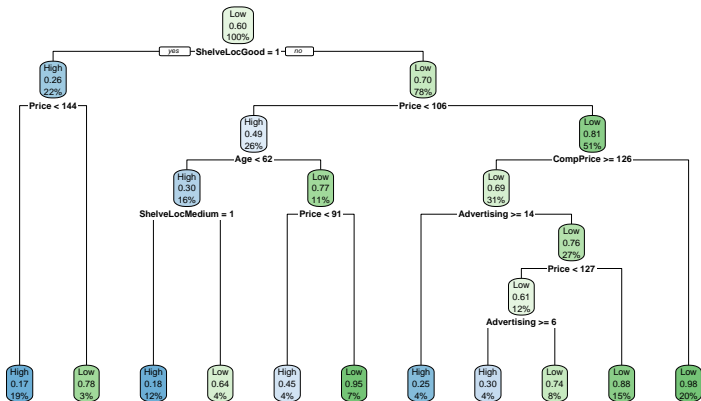
Summary of sample sizes: 252, 252, 251, 252, 252, 253, ...

Resampling results across tuning parameters:


cp	ROC	Sens	Spec
0.000	0.738	0.602	0.746
0.005	0.728	0.587	0.750
0.010	0.713	0.578	0.758
0.015	0.711	0.572	0.764
0.020	0.710	0.572	0.765
0.025	0.706	0.563	0.767

- ▶ The final model

```
rpart.plot(sales_tree$finalModel)
```



```
plot(sales_tree)
```



► Call the confusion matrix

```
cm <- confusionMatrix(data=sales_pred$predicted,  
                      reference=sales_pred$actual,  
                      positive="High")  
cm$table
```

	Reference	
Prediction	High	Low
High	36	11
Low	16	57

```
cm$overall[1]
```

Accuracy
0.775

```
cm$byClass[c(1,2,7)]
```

Sensitivity	Specificity	F1
0.692	0.838	0.727

► ROC curve

```
pred = prediction(sales_pred$p_hat, sales_pred$actual,
                  label.ordering = c("Low", "High"))
roc = performance(pred, "tpr", "fpr")
plot(roc, colorize = T, lwd = 2)
abline(a = 0, b = 1)
auc = performance(pred, measure = "auc")
subtitle = sprintf("AUC: %f", auc@y.values)
mtext(side=3, line=1, at=0, adj=0, cex=0.7, subtitle)
```

