

# NBA 4920/6921 Lecture 16

## Tree Methods: Regression Application

Murat Unal

10/26/2021

```
rm(list=ls())  
options(digits = 3, scipen = 999)  
library(tidyverse)  
library(ISLR)  
library(cowplot)  
library(stargazer)  
library(lmtest)  
library(sandwich)  
library(MASS)  
library(jtools)  
library(caret)  
library(rpart)  
library(rpart.plot)  
library(ROCR)  
set.seed(2)
```

```
Hitters <- ISLR::Hitters  
Hitters <- na.omit(Hitters)
```

► Seperate train and test data

```
train = sample(1:nrow(Hitters), 0.7*nrow(Hitters))
```

# Regression Trees

- ▶ We can fit a regression tree using `rpart` and then visualize it using `rpart.plot`
- ▶ we need to set `method = "anova"`, for classification set `method="class"`
- ▶ The main tuning parameter is `cp`, the complexity parameter

```
hit_tree = rpart(Salary ~ ., data = Hitters[train,],  
                  method="anova")
```

```
hit_tree
```

```
n= 184
```

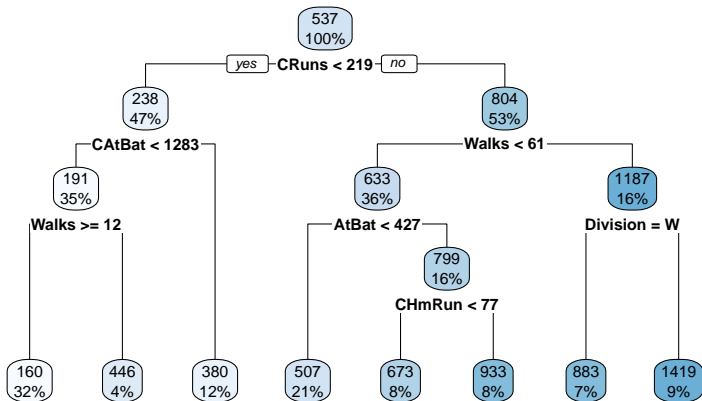
```
node), split, n, deviance, yval
```

```
    * denotes terminal node
```

```
1) root 184 40400000  537
  2) CRuns< 218 87  5140000  238
    4) CAtBat< 1.28e+03 65  4150000  191
      8) Walks>=11.5 58  311000  160 *
      9) Walks< 11.5 7  3330000  446 *
    5) CAtBat>=1.28e+03 22  401000  380 *
  3) CRuns>=218 97 20500000  804
    6) Walks< 61 67  5700000  633
      12) AtBat< 426 38  1720000  507 *
      13) AtBat>=426 29  2570000  799
        26) CHmRun< 76.5 15  335000  673 *
        27) CHmRun>=76.5 14  1750000  933 *
    7) Walks>=61 30  8480000 1190
```

- ▶ We can visualize our tree model with `rpart.plot()`

```
rpart.plot(hit_tree)
```



- ▶ Behind the scenes `rpart()` is automatically applying a range of cost complexity  $\alpha$  values to prune the tree.
- ▶ To compare the error for each  $\alpha$  value, it performs a 10-fold CV (by default)

```
hit_tree$cptable
```

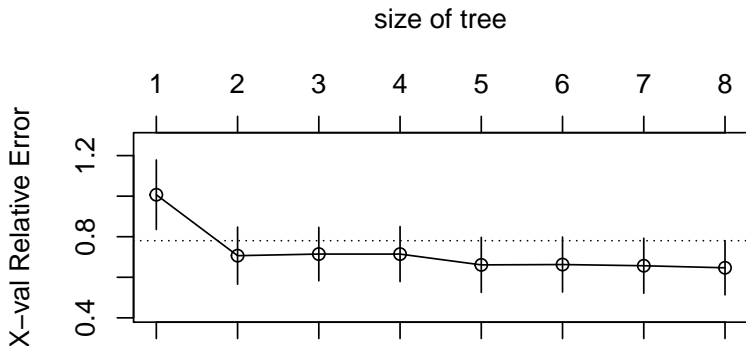
	CP	nsplit	rel error	xerror	xstd
1	0.3639	0	1.000	1.007	0.172
2	0.1576	1	0.636	0.707	0.141
3	0.0524	2	0.479	0.714	0.132
4	0.0348	3	0.426	0.714	0.136
5	0.0145	4	0.391	0.661	0.136
6	0.0126	5	0.377	0.663	0.136
7	0.0122	6	0.364	0.657	0.136
8	0.0100	7	0.352	0.647	0.134

- ▶ Here we don't find much improvement after 2 terminal nodes



- ▶ Notice the dashed line which goes through the point  $|T| = 2$ .
- ▶ It's common to instead use the smallest tree within 1 standard error (SE) of the minimum CV error (this is called the 1-SE rule).
- ▶ Thus, we could use a tree with just 2 terminal nodes and reasonably expect to experience similar results within a small margin of error.


```
plotcp(hit_tree)
```



- ▶ To illustrate the point of selecting a tree with 8 terminal nodes (or 2 if you go by the 1-SE rule), we can force `rpart()` to generate a full tree by setting `cp = 0` (no penalty results in a fully grown tree).

```
hit_tree2 = rpart(Salary ~ ., data = Hitters[train,],  
                  method="anova", control = list(cp=0, xval=10))
```

```
plotcp(hit_tree2)
abline(v = 8, lty = "dashed")
```



- ▶ Make predictions on the test data with best cp

```
pred.tree <- predict(hit_tree, Hitters[-train,])  
  
rmse.min.cp <- sqrt(mean((Hitters[-train,"Salary"]-  
                           pred.tree)^2))  
  
rmse.min.cp
```

```
[1] 309
```

- ▶ `maxdepth`: is the maximum number of internal nodes between the root node and the terminal nodes.
- ▶ We could obtain a tree with 2 terminal nodes by setting `maxdepth=1`.

```
hit_tree1se <- rpart(Salary ~ ., data = Hitters[train,],  
                    method="anova", maxdepth=1)
```

```
hit_tree1se
```

```
n= 184
```

```
node), split, n, deviance, yval  
    * denotes terminal node
```

```
1) root 184 40400000 537
```

```
2) CRuns< 218 87 5140000 238 *
```

```
3) CRuns>=218 97 20500000 804 *
```

► Predictions and rmse

```
pred.tree1se <- predict(hit_tree1se, Hitters[-train,])  
rmse.1se.cp <- sqrt(mean((Hitters[-train,"Salary"]-  
                           pred.tree1se)^2))  
rmse.1se.cp
```

```
[1] 323
```

## Exercise

- ▶ Train a regression tree to predict the crime rates (`crim`) in the Boston dataset.

```
data_test <- read.csv("boston_test.csv")  
data_train <- read.csv("boston_train.csv")
```

► Obtain the tree

```
bos_tree = rpart(crim ~ ., data = data_train,  
                 method="anova")
```

```
bos_tree
```

```
n= 405
```

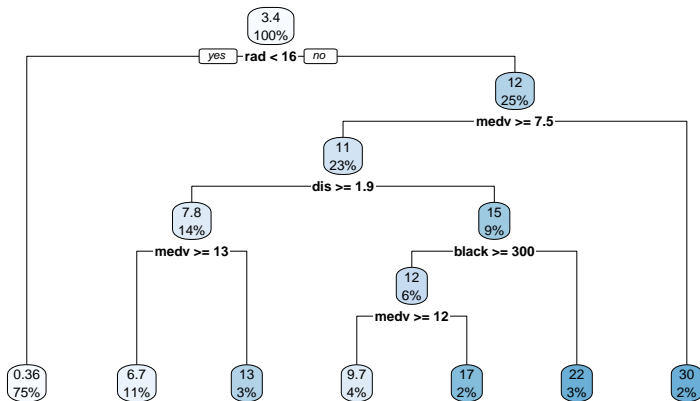
```
node), split, n, deviance, yval  
    * denotes terminal node
```

```
1) root 405 21800.0  3.37  
  2) rad< 16 303    91.0  0.36 *  
  3) rad>=16 102 10800.0 12.30  
    6) medv>=7.45 94  5400.0 10.80  
      12) dis>=1.87 56   750.0  7.82  
        24) medv>=12.6 45   378.0  6.66 *  
        25) medv< 12.6 11    64.3 12.60 *  
      13) dis< 1.87 38  3410.0 15.20  
        26) black>=300 26  1010 0 12 30
```



► Visualize the tree

```
rpart.plot(bos_tree)
```




- ▶ Examine the cross validation errors against cp

```
bos_tree$cptable
```

	CP	nsplit	rel error	xerror	xstd
1	0.5010	0	1.000	1.005	0.252
2	0.1253	1	0.499	0.512	0.176
3	0.0571	2	0.374	0.495	0.157
4	0.0337	3	0.317	0.516	0.170
5	0.0150	4	0.283	0.467	0.145
6	0.0141	5	0.268	0.431	0.133
7	0.0100	6	0.254	0.423	0.133

► Plot cp

```
plotcp(bos_tree)
```



► Obtain the 1se tree

```
bos_tree1se = rpart(crim ~ ., data = data_train,  
                    method="anova", maxdepth=1)  
bos_tree1se
```

n= 405

node), split, n, deviance, yval  
\* denotes terminal node

- 1) root 405 21800 3.37
- 2) rad< 16 303 91 0.36 \*
- 3) rad>=16 102 10800 12.30 \*

► Predictions and rmse for min cp and 1se cp

```
pred.bos.tree <- predict(bos_tree, data_test)
rmse.min.cp <- sqrt(mean((data_test$crim-
                          pred.bos.tree)^2))
rmse.min.cp
```

```
[1] 9.48
```

```
pred.bos.tree1se <- predict(bos_tree1se, data_test)
rmse.1se.cp <- sqrt(mean((data_test$crim-
                          pred.bos.tree1se)^2))
rmse.1se.cp
```

```
[1] 10.7
```