

NBA 4920/6921 Lecture 20

Ensemble Methods: Boosting Application

Murat Unal

11/09/2021

```
rm(list=ls())
options(digits = 3, scipen = 999)
library(tidyverse)
library(ISLR)
library(jtools)
library(caret)
library(glmnet)
library(rpart)
library(rpart.plot)
library(ROCR)
library(ipred)
library(vip)
library(randomForest)
library(gbm)
library(ranger)
set.seed(2)
```

```
Hitters <- ISLR::Hitters  
Hitters <- na.omit(Hitters)  
train = sample(1:nrow(Hitters), 0.7*nrow(Hitters))
```

Boosting


- ▶ Let's run a basic GBM model

```
hit_gbm <- gbm(  
  formula = Salary ~ .,  
  data = Hitters[train,],  
  distribution = "gaussian", # SSE loss function  
  n.trees = 1000,  
  shrinkage = 0.001, #learning rate  
  cv.folds = 10,  
  interaction.depth = 3 #depth of each tree  
)  
# find index for number trees with minimum CV error  
best <- which.min(hit_gbm$cv.error)  
# get MSE and compute RMSE  
sqrt(hit_gbm$cv.error[best])
```

[1] 352

- ▶ Results show cross-validated RMSE of `rsqrt(hit_gbm$cv.error[best])` which we achieved with 1000 trees.
- ▶ Training and cross-validated MSE as trees are added to the GBM algorithm
- ▶ The small learning rate is resulting in very small incremental improvements which means many trees are required

```
gbm.perf(hit_gbm, method = "cv")
```




- ▶ Let's increase the learning rate to take larger steps down the gradient descent

```
hit_gbm2 <- gbm(  
  formula = Salary ~ .,  
  data = Hitters[train,],  
  distribution = "gaussian", # SSE loss function  
  n.trees = 1000,  
  shrinkage = 0.1, #learning rate  
  cv.folds = 10,  
  interaction.depth = 3 #depth of each tree  
)  
# find index for number trees with minimum CV error  
best <- which.min(hit_gbm2$cv.error)  
# get MSE and compute RMSE  
sqrt(hit_gbm2$cv.error[best])
```

[1] 297

```
gbm.perf(hit_gbm2, method = "cv")
```



```
[1] 114
```

- ▶ Make predictions on the test data
- ▶ Like most models, we simply use the predict function;
however, we also need to supply the number of trees to use

```
pred.gbm <- predict.gbm(hit_gbm2,n.trees=1000,Hitters[-train,])  
rmse.gbm <- sqrt(mean((Hitters[-train,"Salary"] -  
                        pred.gbm)^2))  
rmse.gbm
```

```
[1] 319
```


- ▶ A better option than manually tweaking hyperparameters one at a time is to perform a grid search which iterates over every combination of hyperparameter values and allows us to assess which combination tends to perform well.
- ▶ Let's search across 16 models with varying learning rates and tree depth. Let's also vary the minimum number of observations allowed in the trees terminal nodes `n.minobsinnode` and introduce stochastic gradient descent by allowing `bag.fraction < 1`

```
# create hyperparameter grid
hyper_grid <- expand.grid(
  shrinkage = c(.01, .1),
  interaction.depth = c(1, 3),
  n.minobsinnode = c(5, 10),
  bag.fraction = c(.7, .8),
  optimal_trees = 0,
  min_RMSE = 0
)
```

```
# grid search
for(i in 1:nrow(hyper_grid)) {
  print(i)
  # train model
  gbm.tune <- gbm(
    formula = Salary ~ .,
    distribution = "gaussian",
    data = Hitters[train,],
    n.trees = 1000,
    interaction.depth = hyper_grid$interaction.depth[i],
    shrinkage = hyper_grid$shrinkage[i],
    n.minobsinnode = hyper_grid$n.minobsinnode[i],
    bag.fraction = hyper_grid$bag.fraction[i],
    cv.folds = 10)

  # add min training error and trees to grid
  hyper_grid$optimal_trees[i] <- which.min(gbm.tune$cv.error)
  hyper_grid$min_RMSE[i] <- sqrt(min(gbm.tune$cv.error))
}
```

```
hyper_grid %>%  
  arrange(min_RMSE) %>%  
  head(10)
```

shrinkage	interaction	depth	minobs	minoblag	fractionoptimal_trees	min_RMSE
0.01		3	5	0.8	738	286
0.01		3	5	0.7	807	294
0.01		3	10	0.7	775	296
0.10		3	10	0.8	61	298
0.01		3	10	0.8	579	298
0.10		3	5	0.8	89	302
0.10		3	10	0.7	51	305
0.10		3	5	0.7	98	306
0.10		1	10	0.7	135	308
0.01		1	10	0.7	980	315

- Once we have found our top model we train a model with those specific parameters.

```
best.model <- hyper_grid %>%  
  arrange(min_RMSE) %>%  
  head(1)  
best.model
```

shrinkage	interaction	depth	minobs	minnode	bag.fraction	optimal_trees	min_RMSE
0.01		3		5	0.8	738	286

- ▶ Let's re-run the GBM model with optimal hyper parameters

```
hit_gbm.final <- gbm(  
  formula = Salary ~ .,  
  data = Hitters[train,],  
  distribution = "gaussian",  
  n.trees = 1000,  
  interaction.depth = best.model$interaction.depth,  
  shrinkage = best.model$shrinkage,  
  n.minobsinnode = best.model$n.minobsinnode,  
  bag.fraction = best.model$bag.fraction,  
  cv.folds = 10)  
# find index for number trees with minimum CV error  
best <- which.min(hit_gbm.final$cv.error)  
# get MSE and compute RMSE  
sqrt(hit_gbm.final$cv.error[best])
```

[1] 298


- Make predictions on the test data

```
pred.gbm.final <- predict.gbm(hit_gbm.final, n.trees=1000,  
rmse.gbm.final <- sqrt(mean((Hitters[-train,"Salary"] -  
                                pred.gbm.final)^2))  
rmse.gbm.final
```

```
[1] 276
```

► Variable Importance Plot

```
vip(hit_gbm.final)
```



► Compare prediction performance

```
rmse.tree <- 323  
rmse.bag <- 269  
rmse.rf <- 255  
rmse.gbm
```

```
[1] 319
```

```
rmse.gbm.final
```

```
[1] 276
```


Exercise

Boosting

```
cars_train <- read.csv("cayugacars_train.csv")
cars_test  <- read.csv("cayugacars_test.csv")
cars_train$customer_bid <- ifelse(cars_train$customer_bid==0, 1, 0)
cars_test$customer_bid  <- ifelse(cars_test$customer_bid==0, 1, 0)
```


► Run a simple boosting model

```
cars.gbm <- gbm(  
  formula = customer_bid ~ .,  
  data = cars_train,  
  distribution = "bernoulli",  
  n.trees = 1000,  
  shrinkage = 0.1, #learning rate  
  cv.folds = 10,  
  interaction.depth = 3 #depth of each tree  
)  
# find index for number trees with minimum CV error  
best <- which.min(cars.gbm$cv.error)  
# get MSE and compute RMSE  
cars.gbm$cv.error[best]
```

[1] 0.0788

► Plot the cv.error

```
gbm.perf(cars.gbm, method = "cv")
```



[1] 996


► Change the learning rate

```
cars.gbm <- gbm(  
  formula = customer_bid ~ .,  
  data = cars_train,  
  distribution = "bernoulli",  
  n.trees = 1000,  
  shrinkage = 0.2, #learning rate  
  cv.folds = 10,  
  interaction.depth = 3 #depth of each tree  
)  
# find index for number trees with minimum CV error  
best <- which.min(cars.gbm$cv.error)  
# get MSE and compute RMSE  
cars.gbm$cv.error[best]
```

[1] 0.0742

► Plot the cv.error

```
gbm.perf(cars.gbm, method = "cv")
```



[1] 665

► Create hyperparameter grid

```
hyper_grid <- expand.grid(  
  shrinkage = c(.01, .2),  
  interaction.depth = c(1, 3),  
  n.minobsinnode = c(5, 10),  
  bag.fraction = c(.7, .8),  
  optimal_trees = 0,  
  min_RMSE = 0  
)  
# total number of combinations  
nrow(hyper_grid)
```

```
[1] 16
```

► Run the model

```
# grid search
for(i in 1:nrow(hyper_grid)) {
  print(i)
  # train model
  gbm.tune <- gbm(
    formula = customer_bid ~ .,
    distribution = "bernoulli",
    data = cars_train,
    n.trees = 1000,
    interaction.depth = hyper_grid$interaction.depth[i],
    shrinkage = hyper_grid$shrinkage[i],
    n.minobsinnode = hyper_grid$n.minobsinnode[i],
    bag.fraction = hyper_grid$bag.fraction[i],
    cv.folds = 10)
  # add min training error and trees to grid
  hyper_grid$optimal_trees[i] <- which.min(gbm.tune$cv.error)
  hyper_grid$min_RMSE[i] <- sqrt(min(gbm.tune$cv.error))
}
```

► Sort the results

```
hyper_grid %>%  
  arrange(min_RMSE) %>%  
  head(10)
```

shrinkage	interaction	depth	min_obs	min_node	lag	fraction	optimal_trees	min_RMSE
0.20		3		5		0.8	839	0.244
0.20		3		10		0.8	652	0.267
0.20		3		10		0.7	664	0.271
0.20		3		5		0.7	571	0.298
0.01		3		10		0.8	1000	0.551
0.01		3		10		0.7	1000	0.560
0.01		3		5		0.8	1000	0.565
0.01		3		5		0.7	1000	0.565
0.20		1		10		0.7	66	0.918
0.20		1		10		0.8	85	0.921

- ▶ Train a model with the optimal parameters.

```
best.model <- hyper_grid %>%  
  arrange(min_RMSE) %>%  
  head(1)  
best.model
```

shrinkage	interaction	depth	minobs	minnode	lag	fraction	optimal_trees	min_RMSE
0.2		3		5		0.8	839	0.244

- Re-run the GBM model with optimal hyper parameters

```
cars.gbm.final <- gbm(  
  formula = customer_bid ~ .,  
  distribution = "bernoulli",  
  data = cars_train,  
  n.trees = 1000,  
  interaction.depth = best.model$interaction.depth,  
  shrinkage = best.model$shrinkage,  
  n.minobsinnode = best.model$n.minobsinnode,  
  bag.fraction = best.model$bag.fraction,  
  cv.folds = 10)  
# find index for number trees with minimum CV error  
best <- which.min(cars.gbm.final$cv.error)  
cars.gbm.final$cv.error[best]
```

```
[1] 0.0682
```

- Make predictions on the test data and classify into classes

```
pred.gbm.final <- predict.gbm(cars.gbm.final, n.trees=1000,  
                              type="response", cars_test)  
yhat.gbm.final <- as.factor(ifelse(pred.gbm.final>=0.5,1,0))
```

► Confusion Matrix

```
cm <- confusionMatrix(data=yhat.gbm.final,  
reference=as.factor(cars_test$customer_bid),  
positive="1")  
cm$table
```

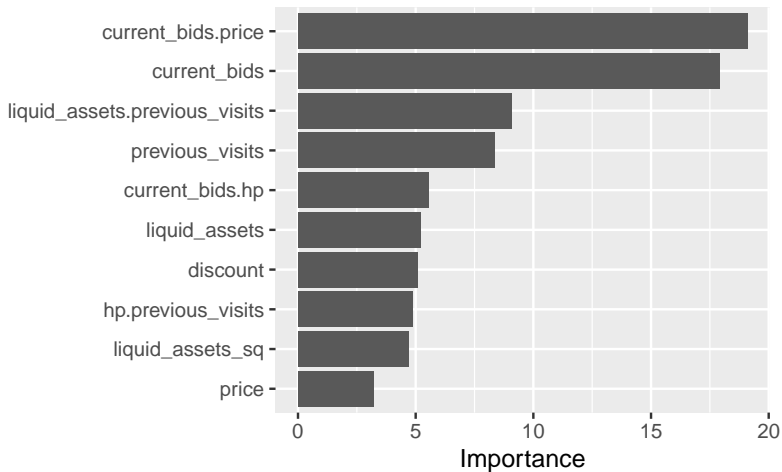
	Reference	
Prediction	0	1
0	187	0
1	1	119

```
c(cm$overall[1],cm$byClass[c(1,2,7)])
```

Accuracy	Sensitivity	Specificity	F1
0.997	1.000	0.995	0.996

► Variable Importance Plot

```
vip(cars.gbm.final)
```



Random Forest

- Define tuning grid

```
hyper_grid <- expand.grid(  
  mtry      = seq(2, 12, by = 2),  
  node_size = seq(2, 8, by = 2),  
  sample_size = c(.5, .70, .80),  
  OOB_RMSE   = 0  
)
```

_ Apply tuning

```
for(i in 1:nrow(hyper_grid)) {  
  # train model  
  model <- ranger(  
    formula      = customer_bid ~ .,  
    data         = cars_train,  
    num.trees    = 1000,  
    mtry         = hyper_grid$mtry[i],  
    min.node.size = hyper_grid$node_size[i],  
    sample.fraction = hyper_grid$sample_size[i]  )  
  # add OOB error to grid  
  hyper_grid$OOB_RMSE[i] <- model$prediction.error  
}
```

► Show tuning results

```
hyper_grid %>%  
  arrange(OOB_RMSE) %>% head(10)
```

mtry	node_size	sample_size	OOB_RMSE
12	4	0.8	0.066
12	2	0.8	0.066
12	6	0.8	0.067
12	8	0.8	0.068
12	2	0.7	0.068
12	4	0.7	0.068
10	2	0.8	0.068
12	6	0.7	0.068
10	4	0.8	0.069
10	6	0.8	0.069

```
best.rf <-hyper_grid %>%  
  arrange(OOB_RMSE) %>% head(1)
```



```
optimal_rf <- ranger(  
  formula      = customer_bid ~ .,  
  data         = cars_train,  
  num.trees    = 1000,  
  mtry         = best.rf$mtry,  
  min.node.size = best.rf$node_size,  
  sample.fraction = best.rf$sample_size,  
  importance    = 'impurity')
```

► Make predictions

```
predict_rf <- predict(optimal_rf, cars_test, type="response")  
y.hat_rf <- ifelse(predict_rf >= 0.5, 1, 0)
```

► Confusion Matrix

```
cm <- confusionMatrix(data=as.factor(y.hat_rf),  
reference=as.factor(cars_test$customer_bid),  
positive="1")  
cm$table
```

	Reference	
Prediction	0	1
0	184	15
1	4	104

```
c(cm$overall[1],cm$byClass[c(1,2,7)])
```

Accuracy	Sensitivity	Specificity	F1
0.938	0.874	0.979	0.916

Bagging

```
cars_bag <- bagging(as.factor(customer_bid) ~ ., data = cars,
                   nbagg=500, coob=TRUE,
                   control=rpart.control(cp=0))
cars_bag
```

Bagging classification trees with 500 bootstrap replication

```
Call: bagging.data.frame(formula = as.factor(customer_bid),
                         nbagg = 500, coob = TRUE, control = rpart.control(cp =
```

Out-of-bag estimate of misclassification error: 0.0833

► Call the confusion matrix

```
cm_bag <- confusionMatrix(data=as.factor(cars_pred_bag$pred),  
                           reference=as.factor(cars_pred_bag$actual),  
                           positive="1")  
cm_bag$table
```

	Reference	
Prediction	0	1
0	182	15
1	6	104

```
c(cm_bag$overall[1],cm_bag$byClass[c(1,2,7)])
```

Accuracy	Sensitivity	Specificity	F1
0.932	0.874	0.968	0.908